



elasticsearch

tutorialspoint

SIMPLY EASY LEARNING

www.tutorialspoint.com



<https://www.facebook.com/tutorialspointindia>



<https://twitter.com/tutorialspoint>

About the Tutorial

Elasticsearch is a real-time distributed and open source full-text search and analytics engine. It is used in Single Page Application (SPA) projects. Elasticsearch is open source developed in Java and used by many big organizations around the world. It is licensed under the Apache license version 2.0.

In this brief tutorial, we will be explaining the basics of Elasticsearch and its features.

Audience

This tutorial is designed for software professionals who want to learn the basics of Elasticsearch and its programming concepts in simple and easy steps. It describes the components of Elasticsearch with suitable examples.

Prerequisites

You should have a basic understanding of Java, JSON, search engines, and web technologies. The interaction with Elasticsearch is through RESTful API; therefore, it is always recommended to have knowledge of RESTful API.

Disclaimer & Copyright

© Copyright 2016 by Tutorials Point (I) Pvt. Ltd.

All the content and graphics published in this e-book are the property of Tutorials Point (I) Pvt. Ltd. The user of this e-book is prohibited to reuse, retain, copy, distribute or republish any contents or a part of contents of this e-book in any manner without written consent of the publisher.

We strive to update the contents of our website and tutorials as timely and as precisely as possible, however, the contents may contain inaccuracies or errors. Tutorials Point (I) Pvt. Ltd. provides no guarantee regarding the accuracy, timeliness or completeness of our website or its contents including this tutorial. If you discover any errors on our website or in this tutorial, please notify us at contact@tutorialspoint.com

Table of Contents

About the Tutorial.....	i
Audience	i
Prerequisites	i
Disclaimer & Copyright.....	i
Table of Contents	ii
1. ELASTICSEARCH — BASIC CONCEPTS.....	1
Elasticsearch – General Features	1
Elasticsearch – Key Concepts	1
Elasticsearch – Advantages	2
Elasticsearch – Disadvantages	2
Comparison between Elasticsearch and RDBMS.....	3
2. ELASTICSEARCH — INSTALLATION	4
3. ELASTICSEARCH – POPULATE ELASTICSEARCH.....	8
Create Index.....	8
Create Mapping and Add data.....	8
Add another Index	9
Create Mapping and Add Data	9
4. ELASTICSEARCH — MIGRATION BETWEEN VERSIONS	11
Snapshot and Restore Module	11
Full Cluster Restart	12
Rolling Upgrades	14

5. ELASTICSEARCH — API CONVENTIONS	15
Multiple Indices	15
Comma Separated Notation	15
_all keyword for all indices.....	15
Wildcards (* , + , -)	16
allow_no_indices	18
expand_wildcards	19
Date Math Support in Index Names	19
Pretty Results.....	20
Human Readable Output.....	21
Response Filtering	21
6. ELASTICSEARCH — DOCUMENT APIS.....	22
Index API.....	22
Automatic Index Creation	22
Versioning	23
Operation Type	23
Automatic ID generation	24
Parents and Children.....	24
Timeout.....	25
Get API	25
Delete API	26
Update API.....	27
Multi Get API.....	28
Bulk API.....	28

7. ELASTICSEARCH — SEARCH APIS	29
Multi-Index	29
Multi-Type	29
URI Search.....	30
Request Body Search.....	30
8. ELASTICSEARCH — AGGREGATIONS	32
Metrics Aggregations	32
Avg Aggregation	32
Cardinality Aggregation.....	33
Extended Stats Aggregation	34
Max Aggregation	35
Min Aggregation	35
Sum Aggregation	36
Bucket Aggregations	36
Aggregation Metadata	37
9. ELASTICSEARCH — INDEX APIS	39
Create Index.....	39
Delete Index.....	41
Get Index	41
Index Exist.....	41
Open / Close Index API.....	42
Index Aliases	42
Index Settings.....	43
Analyze	43
Index Templates.....	44

Index Stats	44
Flush	45
Refresh.....	45
10. ELASTICSEARCH — CLUSTER APIS.....	46
Cluster Health	46
Cluster State.....	46
Cluster Stats	47
Pending Cluster Tasks.....	47
Cluster Reroute	47
Cluster Update Settings.....	48
Node Stats.....	48
Nodes hot_threads.....	49
11. ELASTICSEARCH — QUERY DSL.....	50
Match All Query	50
Full Text Queries	51
Match query.....	51
multi_match query.....	52
Query String Query	52
Term Level Queries	53
Range Query	53
Compound Queries	55
Joining Queries.....	56
Geo Queries	57
12. ELASTICSEARCH — MAPPING	59
Field Data Types.....	59

Mapping Types.....	61
Dynamic Mapping	61
Mapping Parameters.....	62
13. ELASTICSEARCH — ANALYSIS.....	63
Analyzers.....	64
Tokenizers.....	65
Token Filters.....	66
Character Filters	66
14. ELASTICSEARCH — MODULES.....	67
Cluster-Level Routing and Shard Allocation.....	67
Discovery	69
Gateway.....	70
HTTP.....	70
Indices.....	71
Node	72
15. ELASTICSEARCH — TESTING	73
Unit Testing.....	73
Integration Testing.....	73
Randomized Testing	75
Assertions	75

1. Elasticsearch — Basic Concepts

Elasticsearch is an Apache Lucene-based search server. It was developed by Shay Banon and published in 2010. It is now maintained by Elasticsearch BV. Its latest version is 2.1.0.

Elasticsearch is a real-time distributed and open source full-text search and analytics engine. It is accessible from RESTful web service interface and uses schema less JSON (JavaScript Object Notation) documents to store data. It is built on Java programming language, which enables Elasticsearch to run on different platforms. It enables users to explore very large amount of data at very high speed.

Elasticsearch – General Features

The general features of Elasticsearch are as follows:

- Elasticsearch is scalable up to petabytes of structured and unstructured data.
- Elasticsearch can be used as a replacement of document stores like MongoDB and RavenDB.
- Elasticsearch uses denormalization to improve the search performance.
- Elasticsearch is one of the popular enterprise search engines, which is currently being used by many big organizations like Wikipedia, The Guardian, StakOverflow, GitHub etc.
- Elasticsearch is open source and available under the Apache license version 2.0.

Elasticsearch – Key Concepts

The key concepts of Elasticsearch are as follows:

- **Node:** It refers to a single running instance of Elasticsearch. Single physical and virtual server accommodates multiple nodes depending upon the capabilities of their physical resources like RAM, storage and processing power.
- **Cluster:** It is a collection of one or more nodes. Cluster provides collective indexing and search capabilities across all the nodes for entire data.

- **Index:** It is a collection of different type of documents and document properties. Index also uses the concept of shards to improve the performance. For example, a set of document contains data of a social networking application.
- **Type/Mapping:** It is a collection of documents sharing a set of common fields present in the same index. For example, an Index contains data of a social networking application, and then there can be a specific type for user profile data, another type for messaging data and another for comments data.
- **Document:** It is a collection of fields in a specific manner defined in JSON format. Every document belongs to a type and resides inside an index. Every document is associated with a unique identifier, called the UID.
- **Shard:** Indexes are horizontally subdivided into shards. This means each shard contains all the properties of document, but contains less number of JSON objects than index. The horizontal separation makes shard an independent node, which can be store in any node. Primary shard is the original horizontal part of an index and then these primary shards are replicated into replica shards.
- **Replicas:** Elasticsearch allows a user to create replicas of their indexes and shards. Replication not only helps in increasing the availability of data in case of failure, but also improves the performance of searching by carrying out a parallel search operation in these replicas.

Elasticsearch – Advantages

- Elasticsearch is developed on Java, which makes it compatible on almost every platform.
- Elasticsearch is real time, in other words after one second the added document is searchable in this engine.
- Elasticsearch is distributed, which makes it easy to scale and integrate in any big organization.
- Creating full backups are easy by using the concept of gateway, which is present in Elasticsearch.
- Handling multi-tenancy is very easy in Elasticsearch when compared to Apache Solr.
- Elasticsearch uses JSON objects as responses, which makes it possible to invoke the Elasticsearch server with a large number of different programming languages.

- Elasticsearch supports almost every document type except those that do not support text rendering.

Elasticsearch – Disadvantages

- Elasticsearch does not have multi-language support in terms of handling request and response data (only possible in JSON) unlike in Apache Solr, where it is possible in CSV, XML and JSON formats.
- Elasticsearch also have a problem of Split brain situations, but in rare cases.

Comparison between Elasticsearch and RDBMS

In Elasticsearch, index is a collection of type just as database is a collection of tables in RDBMS (Relation Database Management System). Every table is a collection of rows just as every mapping is a collection of JSON objects Elasticsearch.

Elasticsearch	RDBMS
Index	Database
Shard	Shard
Mapping	Table
Field	Field
JSON Object	Tuple

2. Elasticsearch — Installation

The steps for installation of Elasticsearch are as follows:

Step 1: Check the minimum version of your java in installed your computer, it should be java 7 or more updated version. You can check by doing the following:

In Windows Operating System (OS) (using command prompt):

```
> java -version
```

In UNIX OS (Using Terminal):

```
$ echo $JAVA_HOME
```

Step 2: Download Elasticsearch from <https://www.elastic.co/downloads/elasticsearch>

- For windows OS download ZIP file.
- For UNIX OS download TAR file.
- For Debian OS download DEB file.
- For Red Hat and other Linux distributions download RPN file.
- APT and Yum utilities can also be used to install Elasticsearch in many Linux distributions.

Step 3: Installation process for Elasticsearch is very easy and described below for different OS:

- **Windows OS:** Unzip the zip package and the Elasticsearch is installed.
- **UNIX OS:** Extract tar file in any location and the Elasticsearch is installed.

```
$tar -xvf elasticsearch-2.1.0.tar.gz
```

- **Using APT utility for Linux OS:**
 - Download and install the Public Signing Key:

```
$ wget -qO - https://packages.elastic.co/GPG-KEY-elasticsearch |
sudo apt-key add -
```

- Save the repository definition:

```
$ echo "deb http://packages.elastic.co/elasticsearch/2.x/debian
stable main" | sudo tee -a /etc/apt/sources.list.d/elasticsearch-
2.x.list
```

- Run update:

```
$ sudo apt-get update
```

- Now you can install by using the following command:

```
$ sudo apt-get install elasticsearch
```

- **Using YUM utility for Debian Linux OS:**

- Download and install the Public Signing Key:

```
$ rpm --import https://packages.elastic.co/GPG-KEY-elasticsearch
```

- ADD the below text in the file with .repo suffix in your "/etc/yum.repos.d/" directory. For example, **elasticsearch.repo**

```
[elasticsearch-2.x]
name=Elasticsearch repository for 2.x packages
baseurl=http://packages.elastic.co/elasticsearch/2.x/centos
gpgcheck=1
gpgkey=http://packages.elastic.co/GPG-KEY-elasticsearch
enabled=1
```

- You can now install Elasticsearch by using the following command:

```
$ yum install elasticsearch
```

Step 4: Go to the Elasticsearch home directory and inside the bin folder. Run the elasticsearch.bat file in case of windows or you can do the same using command prompt and through terminal in case of UNIX run Elasticsearch file.

In Windows:

```
> cd elasticsearch-2.1.0/bin  
> elasticsearch
```

In Linux:

```
$ cd elasticsearch-2.1.0/bin  
$ ./elasticsearch
```

Note: in case of windows, you might get error stating JAVA_HOME is not set, please set it in environment variables to "C:\Program Files\Java\jre1.8.0_31" or the location where you installed java.

Step 5: Default port for Elasticsearch web interface is 9200 or you can change it by changing http.port inside elasticsearch.yml file present in bin directory. You can check if the server is up and running by browsing <http://localhost:9200>. It will return a JSON object, which contains the information about the installed Elasticsearch in the following way:

```
{  
  "name" : "Brain-Child",  
  "cluster_name" : "elasticsearch",  
  "version" : {  
    "number" : "2.1.0",  
    "build_hash" : "72cd1f1a3eee09505e036106146dc1949dc5dc87",  
    "build_timestamp" : "2015-11-18T22:40:03Z",  
    "build_snapshot" : false,  
    "lucene_version" : "5.3.1"
```

```

    },
    "tagline" : "You Know, for Search"
  }

```

Step 6: You can install fiddler2 from <http://www.telerik.com/download/fiddler> as a front end for your Elasticsearch.

- In the configure window of fiddler2, you can hit the address of Elasticsearch adding an index and if you want, then the type/mapping also using HTTP POST method, for example:

Address bar

```
http://localhost:9200/schools/school
```

Request body

You can add JSON object, which will get store into that index.

- You can use the same for searching anything by just adding "_search" keyword at the end of URL and sent a query in request body for example:

Address bar

```
POST http://localhost:9200/city/schools/_search
```

Request body

```

{
  "query":{
    "match_all":{}
  }
}

```

This query will return everything from that index, which belongs to that particular type.

- You can delete a particular index or type by just putting the URL of the same in address bar and hit it with HTTP DELETE method.

3. Elasticsearch – Populate Elasticsearch

In this section, we will add some index, mapping and data to Elasticsearch. This data will be used in the examples explained in this tutorial.

Create Index

```
POST http://localhost:9200/schools
```

Request Body

It can contain index specific settings, but for now, it is empty for default settings.

Response

```
{"acknowledged": true}
```

(This means index is created)

Create Mapping and Add data

Elasticsearch will auto-create the mapping according to the data provided in request body, we will use its bulk functionality to add more than one JSON object in this index.

```
POST http://localhost:9200/schools/_bulk
```

Request Body

```
{"index":{"_index":"schools", "_type":"school", "_id":"1"}}
{"name":"Central School", "description":"CBSE Affiliation", "street":"Nagan",
"city":"paprola", "state":"HP", "zip":"176115",
"location":[31.8955385,76.8380405], "fees":2000,
"tags":["Senior Secondary", "beautiful campus"],"rating":"3.5"}

{"index":{"_index":"schools", "_type":"school", "_id":"2"}}
{"name":"Saint Paul School", "description":"ICSE
Afiliation", "street":"Dawarka", "city":"Delhi", "state":"Delhi", "zip":"110075", "loca
tion":[28.5733056,77.0122136], "fees":5000,
```



```
"tags":["Good Faculty", "Great Sports"],"rating":"4.5"}

{"index":{"_index":"schools", "_type":"school", "_id":"3"}}
{"name":"Crescent School", "description":"State Board Affiliation", "street":"Tonk Road", "city":"Jaipur", "state":"RJ",
"zip":"176114", "location":[26.8535922,75.7923988], "fees":2500, "tags":["Well equipped labs"],"rating":"4.5"}
```

Response

```
{"took":328,"errors":false,"items":[{"index":{"_index":"schools","_type":"school",
"_id":"1","_version":1,"_shards":{"total":2,"successful":1,"failed":0},"status":201}},{"index":{"_index":"schools","_type":"school","_id":"2","_version":1,"_shards":
{"total":2,"successful":1,"failed":0},"status":201}},{"index":{"_index":"schools",
"_type":"school","_id":"3","_version":1,"_shards":{"total":2,"successful":1,"fail
ed":0},"status":201}}]}
```

Add another Index

Create Index

```
POST http://localhost:9200/schools_gov
```

Request Body

It can contain index specific settings, but for now it's empty for default settings.

Response

```
{"acknowledged": true} (This means index is created)
```

Create Mapping and Add Data

```
POST http://localhost:9200/schools_gov/_bulk
```

Request Body

```
{
  "index": { "_index": "schools_gov", "_type": "school", "_id": "1" }
},
{
  "name": "Model School", "description": "CBSE Affiliation",
  "street": "silk city", "city": "Hyderabad", "state": "AP", "zip": "500030",
  "location": [17.3903703, 78.4752129], "fees": 200,
  "tags": ["Senior Secondary", "beautiful campus"], "rating": "3"
},
{
  "index": { "_index": "schools_gov", "_type": "school", "_id": "2" }
},
{
  "name": "Government School", "description": "State Board Affiliation",
  "street": "Hinjewadi", "city": "Pune", "state": "MH", "zip": "411057",
  "location": [18.599752, 73.6821995], "fees": 500, "tags": ["Great Sports"], "rating": "4"
}
```

Response

```
{
  "took": 179, "errors": false, "items": [
    {
      "index": {
        "_index": "schools_gov",
        "_type": "school",
        "_id": "1",
        "_version": 1,
        "_shards": { "total": 2, "successful": 1, "failed": 0 },
        "status": 201
      }
    },
    {
      "index": {
        "_index": "schools_gov",
        "_type": "school",
        "_id": "2",
        "_version": 1,
        "_shards": { "total": 2, "successful": 1, "failed": 0 },
        "status": 201
      }
    }
  ]
}
```

End of ebook preview
If you liked what you saw...
Buy it from our store @ **<https://store.tutorialspoint.com>**